

**MACHINE
INTELLIGENCE
COMMUNITY**

Machine Intelligence

Rachel Manzelli, Tyrone Hou, Justin Chen
September 12, 2018



Welcome to BUMIC!


We're excited to start a great year with you



Sign in

[https://goo.gl/72uwx](https://goo.gl/72uwxS)
S

Who are we?

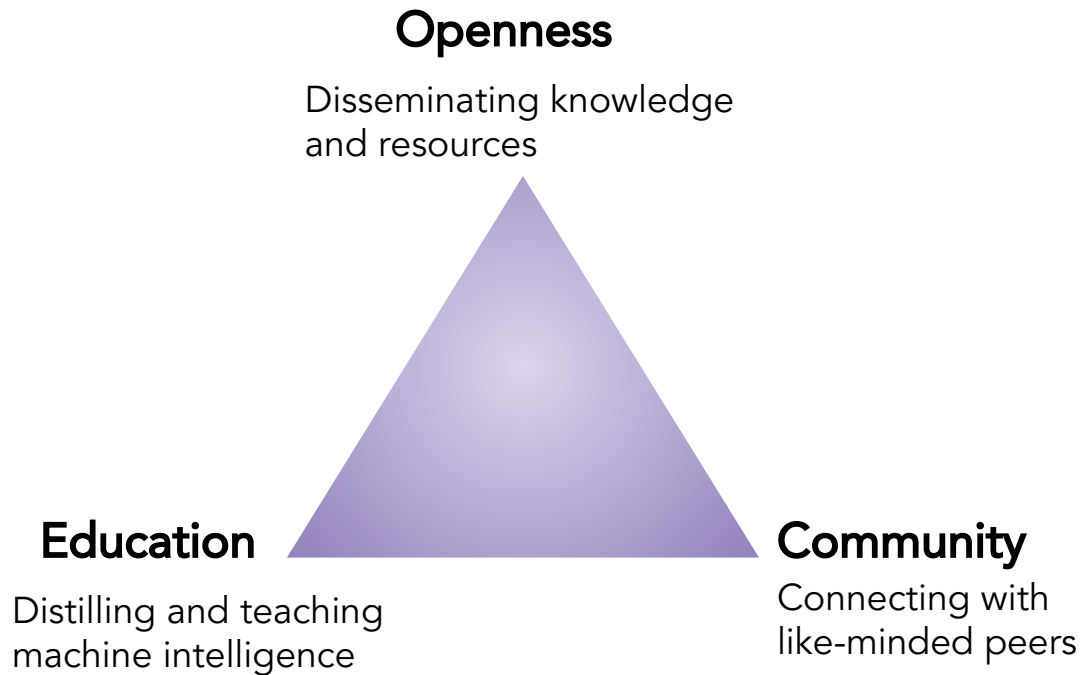


The Machine Intelligence Community is an organization focused on providing opportunities for students to learn about machine intelligence in a community environment.

Who are we?



Our Core Values



Goals of Workshop Series

1. Introduce the field of deep learning
2. Familiarity with terminology, history, concepts, current events, and papers
3. Learn how to read research papers
4. Meet others interested in machine intelligence
5. Develop skill set to start your own deep learning projects



Workshop Topics


1. Machine Intelligence (you are here!)
2. Linear Algebra Foundations in Deep Learning
3. Gradient-Based Learning
4. Neural Networks
5. Regularization
6. Compositional Data
7. Transfer Learning
8. Sequential Data
9. Deep Reinforcement Learning
10. Unsupervised Learning
11. Evolutionary Algorithms




AI, Machine Learning, & Deep Learning

Let's clear some things up...

What is Artificial Intelligence?



<https://www.forbesmiddleeast.com/en/artificial-intelligence-to-add-182-billion-to-u-a-e-s-economy-by-2035/>



<https://afisha.tut.by/news/anews/475339.html?crnd=48542>

What is Artificial Intelligence?



<https://www.forbesmiddleeast.com/insights/ai-intelligence-to-add-182-billion-to-u-a-e-s-economy-by-2035/>



<https://afisha.tut.by/news/anews/478096?id=48542>

What is Artificial Intelligence, really?



Thinking humanly - “[The automation of] activities that we associate with human thinking, activities such as decision-making, problem solving, learning ...” (Bellman, 1978)

Thinking rationally - “The study of the computations that make it possible to perceive, reason, and act.” (Winston, 1992)

Acting humanly - “The study of how to make computers do things at which, at the moment, people are better.” (Rich and Knight, 1991)

Acting rationally - “Computational Intelligence is the study of the design of intelligent agents.” (Poole et al., 1998)

What is Machine Learning?

Coined by Arthur Samuel in 1959 - "A field of study that gives computers the ability to learn without being explicitly programmed."

Applied statistics

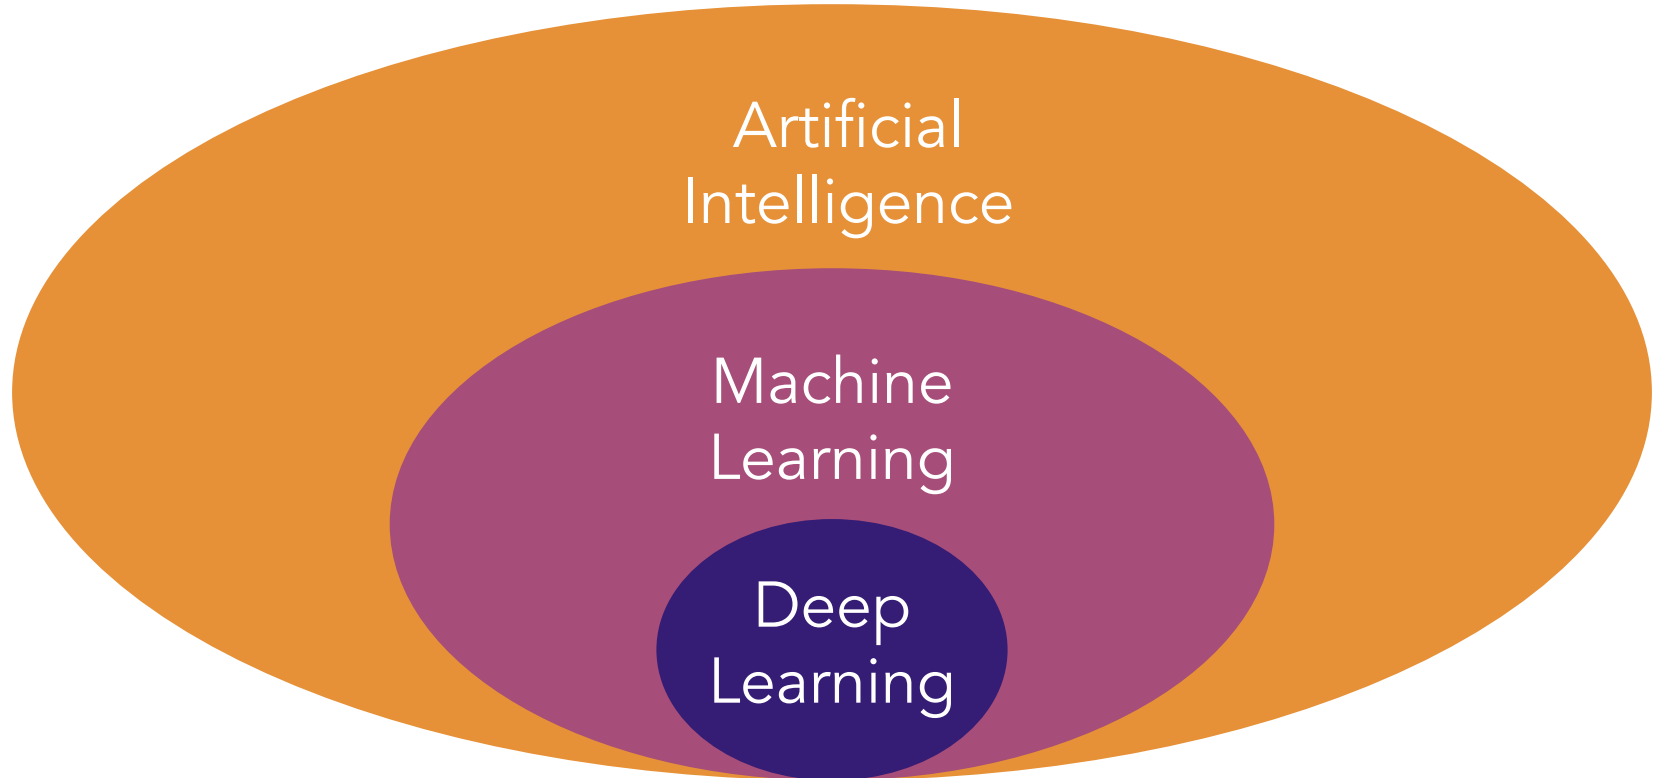
Data-driven

Prediction



<https://history-computer.com/ModernComputer/thinkers/Samuel.html>



What is Deep Learning?



What is Deep Learning?

DEEP NEURAL NETWORK

Input layer → Hidden layer 1 → Hidden layer 2 → Hidden layer 3 → Output layer



<https://medium.freecodecamp.org/an-intuitive-guide-to-convolutional-neural-networks-260c2de0a050>

What is Deep Learning, really?

Calculus

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

Linear Algebra

$$\begin{bmatrix} x_{11} & x_{12} & x_{13} & \dots & x_{1n} \\ x_{21} & x_{22} & x_{23} & \dots & x_{2n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ x_{d1} & x_{d2} & x_{d3} & \dots & x_{dn} \end{bmatrix}$$

Probability & Statistics

$$P(y = j|x) = \frac{e^{x^T w_j}}{\sum_{n=1}^N e^{x^T w_n}}$$

Optimization Theory

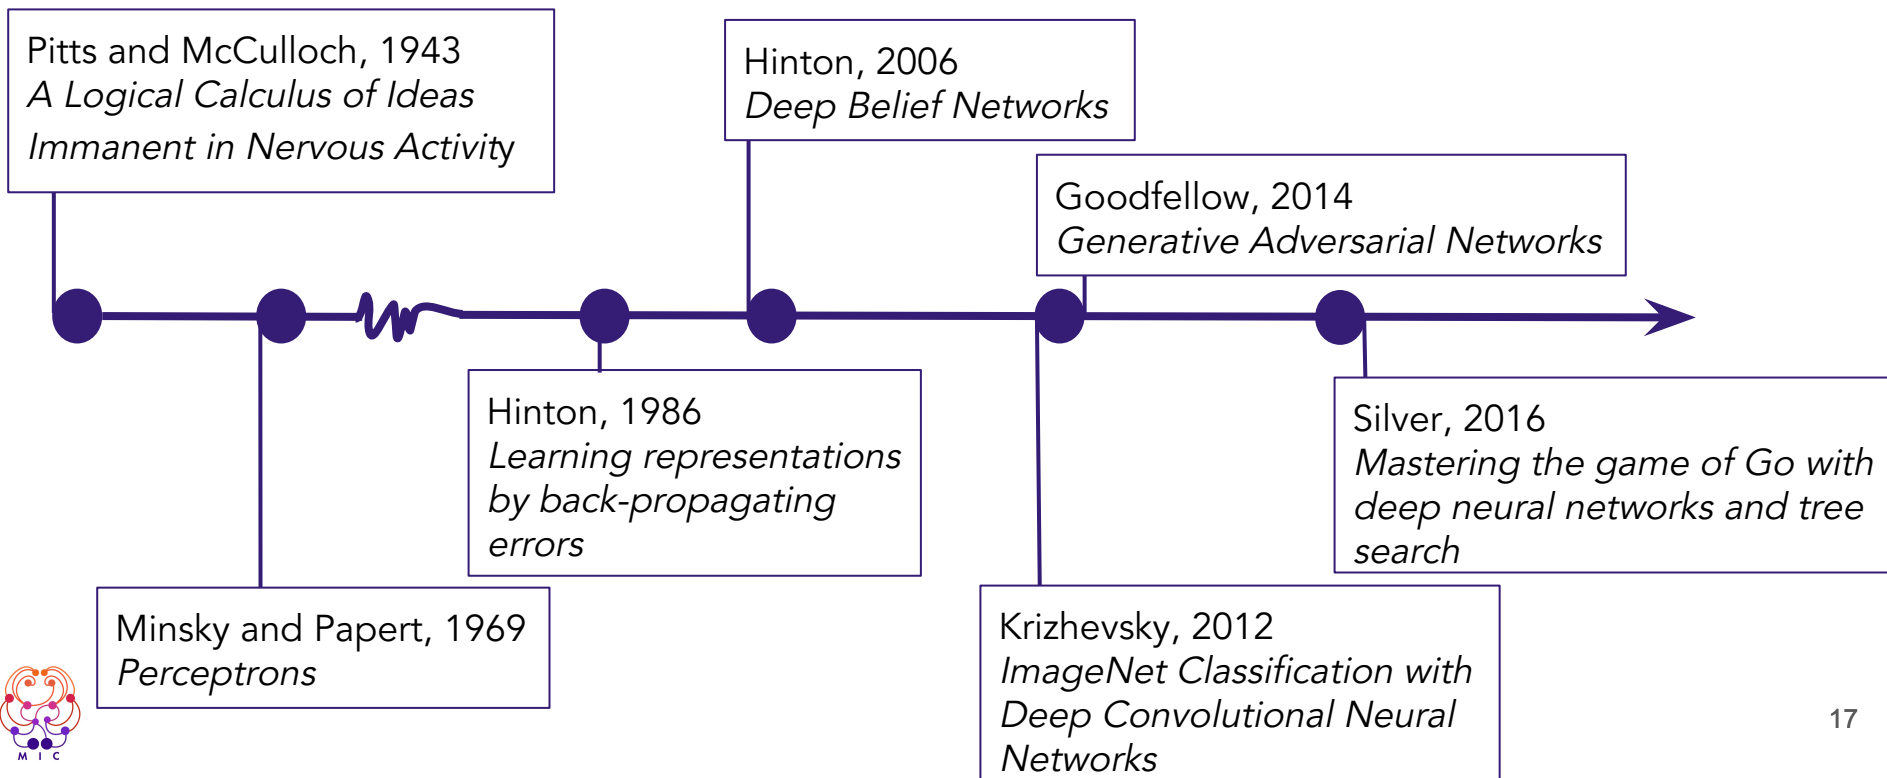
$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Learning Theory


$$Q_{\hat{\pi}}(s_t, a_t) = \mathbb{E}[R_t + \lambda \max_a Q_{\hat{\pi}}(S_{t+1}, a)]$$



Deep Learning Milestones



Rosenblatt's Perceptron



Schematic of Rosenblatt's perceptron.

https://sebastianraschka.com/Articles/2015_singlelayer_neurons.html


Where is Deep Learning?



Microsoft



Tencent 腾讯



SONY

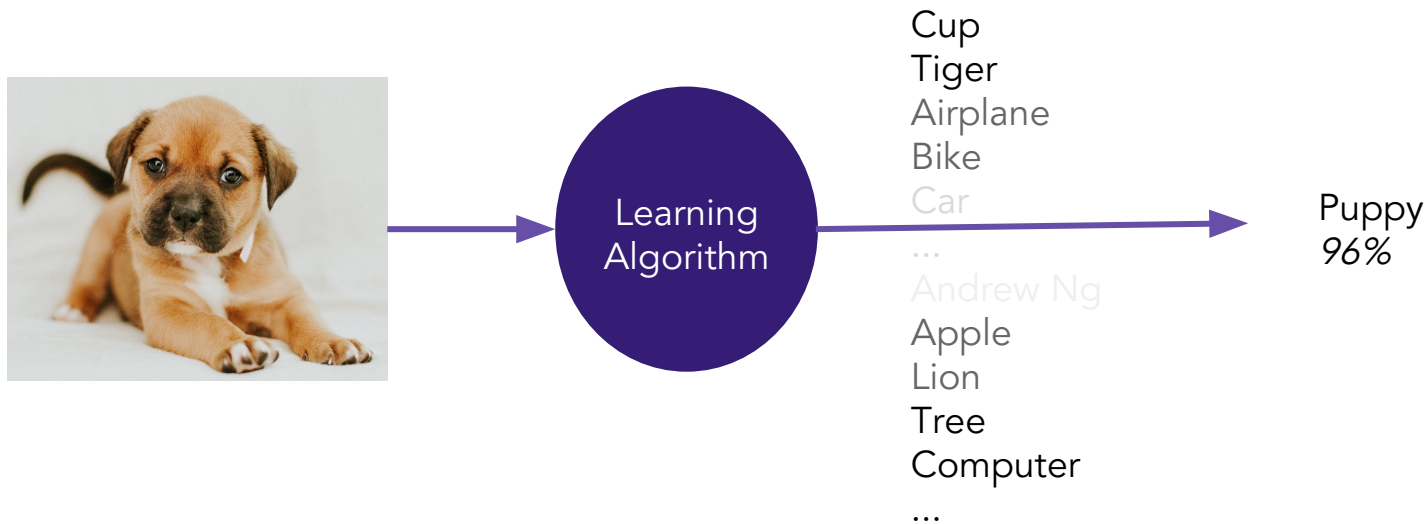




What Can Machines Learn?


Classification, Regression, & Generation

Classification

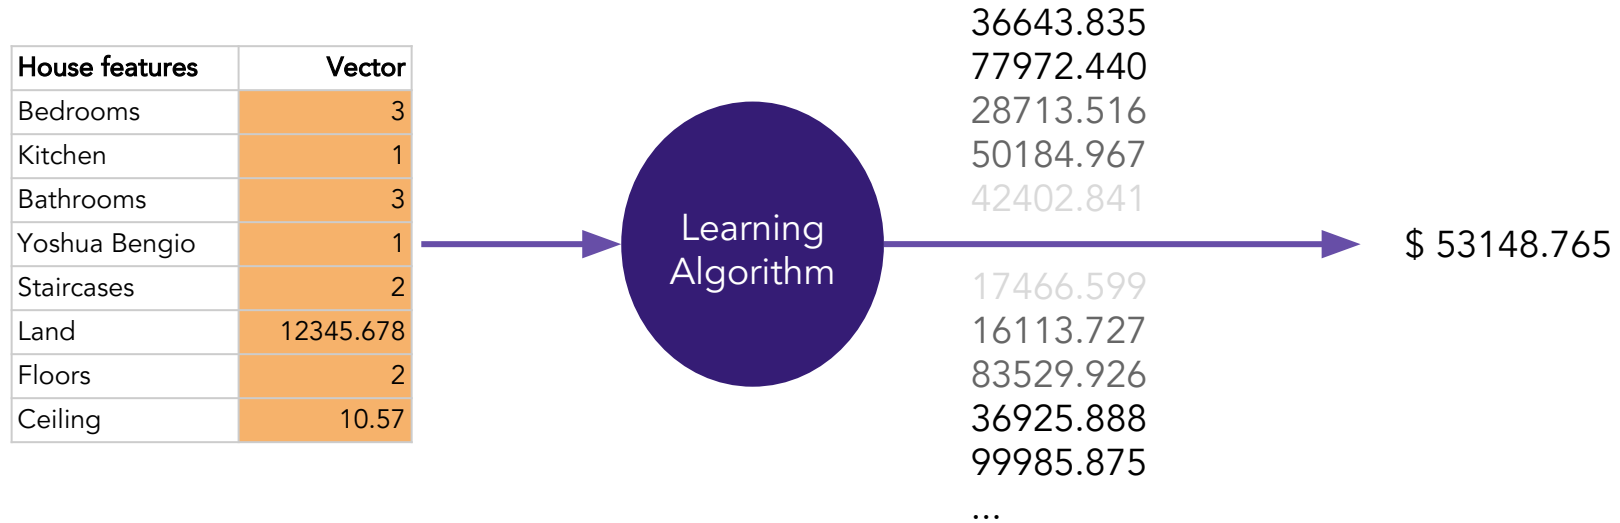


Classification involves predicting classes, where each class is a **discrete value**.

Classification: MNIST




Regression



Regression involves prediction of **continuous real values**.

Generation



*Do you bite your
thumb at us, sir?*

Learning
Algorithm

PANDARUS:

Alas, I think he shall be come
approached and the day When
little strain would be attain'd into
being never fed, And who is but a
chain and subjects of his death, I
should not sleep. Second

Senator:

They are away this miseries,
produced upon my soul, Breaking
and strongly should be buried,
when I perish The earth and
thoughts of many states.

Geoffrey Hinton: Well, your wit is in
the care of side and that.

...

Generative algorithms model and generate the **original distribution of the training data.**

Generation



Keaton Patti ✓
@KeatonPatti

Following



I forced a bot to watch over 1,000 hours of Diners, Drive-Ins, and Dives and then asked it to write a Diners, Drive-Ins, and Dives episode of its own. Here is the first page.



DINERS, DRIVE-INS, AND DIVES

INT. PARKING LOT

GUY FIERI sits in a convertible. He looks like America.

GUY FIERI

I'm Guy Fieri and there's nothing you can do about it. Today I'm eating it all.

Guy takes a bite out of his hair.

INT. DINER'S KITCHEN

Guy and a CHEF stand in a kitchen. Guy has 3 pairs of sunglasses on. The sun can't get him.

GUY FIERI

Prove to me you can panini!

The chef starts boiling a pot of milk. He's scared.

CHEF

Flavortown is near.





How Do Machines Learn?

Learning Algorithms, Data, Training & Testing

Types of Learning

Supervised - labeled dataset


Semi-supervised - partially labeled dataset

Unsupervised - unlabeled dataset

Reinforcement - punishment and reward

Meta - learning to learn

Transfer - use learned parameters from a similar task



Dataset

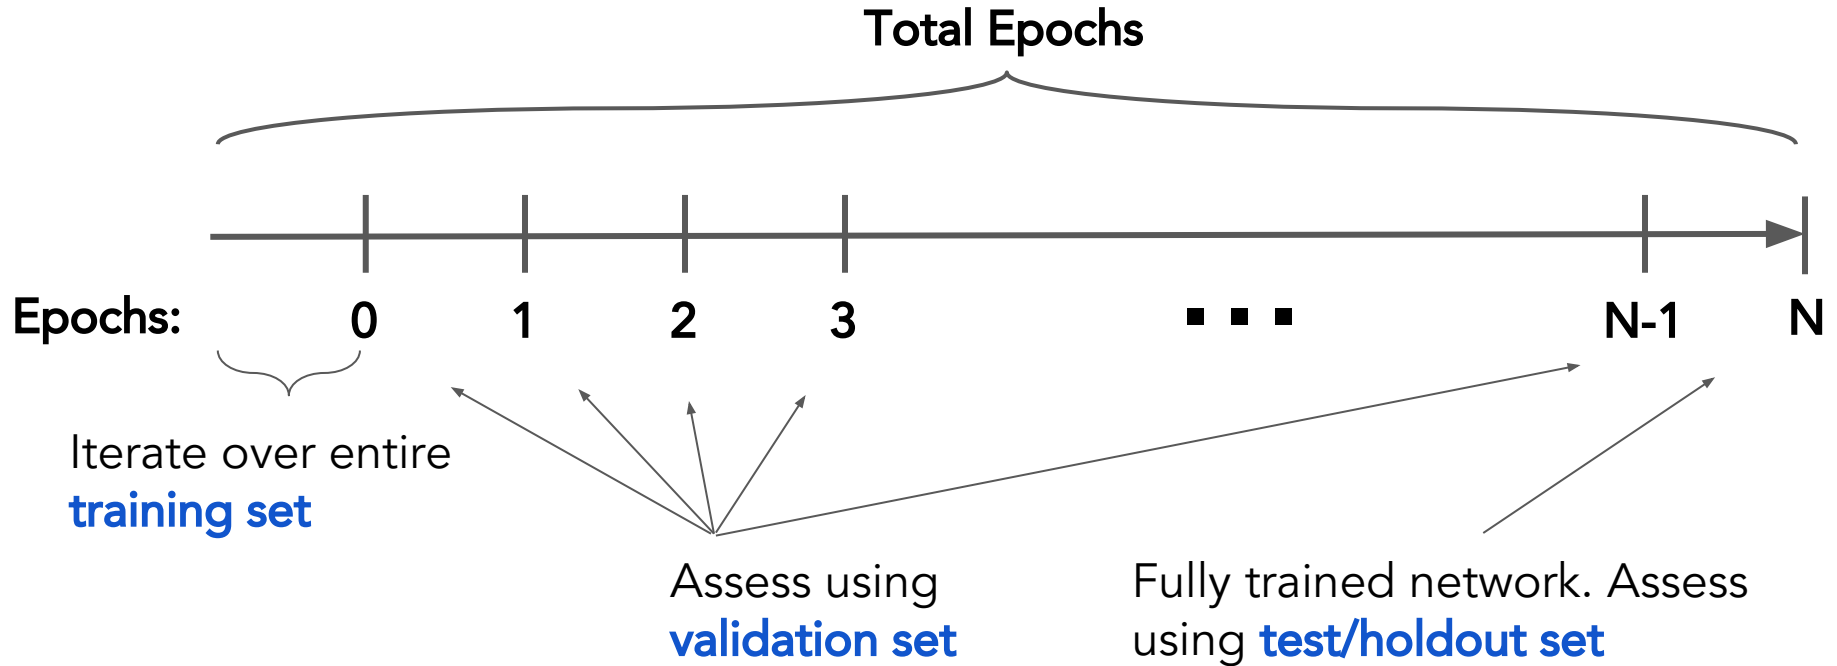
Split - percentages depend on available data



Balance - data distributed close to uniformly



Training, Validation, & Testing



References and Further Reading

1. Russell, Stuart, Peter Norvig, and Artificial Intelligence. "[A modern approach.](#)" Artificial Intelligence. Prentice-Hall, Englewood Cliffs 25 (1995): 27.
2. McCulloch, Warren S., and Walter Pitts. "[A logical calculus of the ideas immanent in nervous activity.](#)" The bulletin of mathematical biophysics 5.4 (1943): 115-133.
3. Ivakhnenko, Alekseï Grigor'evich, and Valentin Grigorévich Lapa. [Cybernetic predicting devices.](#) No. TR-EE66-5. PURDUE UNIV LAFAYETTE IND SCHOOL OF ELECTRICAL ENGINEERING, 1966.
4. LeCun, Yann, et al. "[Gradient-based learning applied to document recognition.](#)" Proceedings of the IEEE 86.11 (1998): 2278-2324.
5. Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "[Imagenet classification with deep convolutional neural networks.](#)" Advances in neural information processing systems. 2012.
6. Halevy, Alon, Peter Norvig, and Fernando Pereira. "[The unreasonable effectiveness of data.](#)" IEEE Intelligent Systems 24.2 (2009): 8-12.

Additional reading topics: General adversarial networks, neural turing machines, defeat of Lee Sedol with AlphaGo



Follow us for cool MI things!



<https://www.facebook.com/BostonUniversityMIC/>

and

<https://www.facebook.com/groups/1017831178363028/>



<https://goo.gl/KHRCxw>



<http://eepurl.com/c2lXqn>

Upcoming Events with BUMIC

- Next workshop: **Linear Algebra Foundations in Deep Learning**
- First Hack Night:
- First paper discussion: TBD



Thank you for joining us!

See you next week...

