




Please sign-in: goo.gl/4R9Zca



Gradient-Based Learning

BOSTON UNIVERSITY
MACHINE INTELLIGENCE
COMMUNITY

Devin de Hueck, Justin Chen
Sept. 19, 2018

Goals of this Workshop


- Introduction to statistical learning
- Cost/error functions
- Gradient descent
- Variations of gradient descent






The learning problem


We have some data D



Make an assumption about D




$$y = b + mx$$



$$\hat{f} = \theta_0 + \theta_1 x$$

What is learning?

The approximation of some unknown function f based on some data D .



$$f : X \rightarrow Y$$

$$\hat{f} = \theta_0 + \theta_1 x$$

How do we set the parameters?
With the use of **gradients**.




What are gradients?

Univariate Functions


Univariate function

$$f(x) = x^2$$



Derivative

$$\frac{\delta f}{\delta x} = 2x$$



Multivariate functions

Multivariate function

Function of more than one variable/**dimension**

$$f(\vec{x}) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 \dots$$

Multivariate derivative

Shows us slope of a function in all dimensions

ex.

$$f(x_0, x_1) = 4x_0 + 8x_1$$

$$\nabla f = [4, 8]$$

$$\nabla f = \begin{bmatrix} \frac{\delta f}{\delta x_0} \\ \frac{\delta f}{\delta x_1} \\ \vdots \\ \frac{\delta f}{\delta x_n} \end{bmatrix}$$



The gradient tells us how much the output of some function will change with respect to a parameter of that function.



How are gradients used to learn?

Cost Functions

- Learning algorithms need to measure how wrong it is in order to improve the model
- A **Cost function** measures the **error** distance between a **hypothesis** and a **label** (correct answer)
- For example
 - Learning algorithm's hypothesis: [0.1, 0.2, 0.3]
 - Correct answers: [0.5, 0.8, 0.2]
 - Errors: [-0.4, -0.6, 0.1]
 - Least Square Error (LSE): $\frac{1}{2}[(0.1-0.5)^2+(0.2-0.8)^2+(0.3-0.2)^2]$



Different Cost Functions

- Common terms: **cost/criterion/objective/loss function**
- Tailored for the model
- Common cost functions:

Regression: **Least Square Error**

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Classification: **Negative Log Likelihood a.k.a Cross Entropy**

$$J(\theta) = \sum_{i=1}^m (y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})))$$

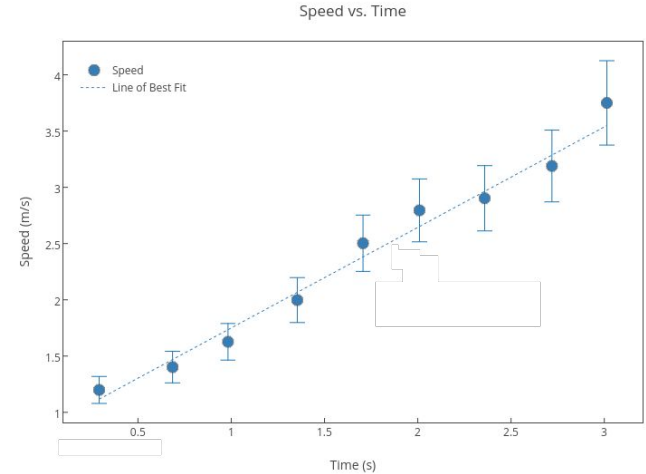


LSE Explained

$h_{\theta}(x^{(i)})$ - Your model's prediction

$y^{(i)}$ - The correct answer (label)

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



<https://plotlyblog.tumblr.com/post/84309369787/best-fit-lines-in-plotly>



Optimizing Cost Function


- Goal is to **minimize cost function** J
 - Compute derivative of J w.r.t. parameters θ

$$\nabla J = \frac{\delta J}{\delta \theta}$$

- Consider this simple cost function

$$f(x) = x^2 \rightarrow \frac{\delta f}{\delta x} = 2x$$

- Solve derivative for 0
- **Convex** functions have single **global minima**
- Most **cost landscapes** are **non-convex** - contain many **local minima** (exponentially many in the number of dimensions)




<https://sebastianraschka.com/faq/docs/closed-form-vs-gd.html>



Gradient Descent

- Assumption:
 - Inputs are sampled i.i.d.
- Use gradient to iteratively traverse parameter landscape
- Gradient is direction of steepest ascent



<https://sebastianraschka.com/faq/docs/closed-form-vs-gd.html>

Gradient Descent

Also written $\nabla J(\theta)$

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

Scalar learning rate

Individual weights

Cost/objective/loss function




Vector of weights

The diagram illustrates the gradient descent update rule. The equation is $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$. Annotations include: 'Scalar learning rate' pointing to α ; 'Individual weights' pointing to the θ_j on the left; 'Cost/objective/loss function' pointing to $J(\theta)$; 'Vector of weights' pointing to θ in the denominator of the derivative; and 'Also written $\nabla J(\theta)$ ' with a bracket pointing to the derivative term $\frac{\partial}{\partial \theta_j} J(\theta)$.




Learning rate

- Gradient points in direct of steepest ascent, but we use the learning rate to control step size
- Typical learning rates to try:
0.1, 0.01, 0.001, 0.0001



Saddle points

- If you are right on the saddle, the gradient does not always help you get off
- Can exist above one dimension
- Many solutions to this problem exist
 - A simple one is by just adding noise, you can force the algorithm to randomly “catch on” to the downward slope




<http://cs231n.github.io/neural-networks-3/>



Different types of Gradient Descent

- Batch/Vanilla gradient descent (which we have been describing so far)
- Stochastic gradient descent
- A third variety is mini-batch gradient descent
 - Between BGD and SGD
- Adagrad, Adam



<http://blog.hackerearth.com/3-types-gradient-descent-algorithms-small-large-data-sets>


Batch/Vanilla Gradient Descent (GD)

Batch size of the entire dataset. Descend in the steepest direction given information from every training example in the dataset.

Loop {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

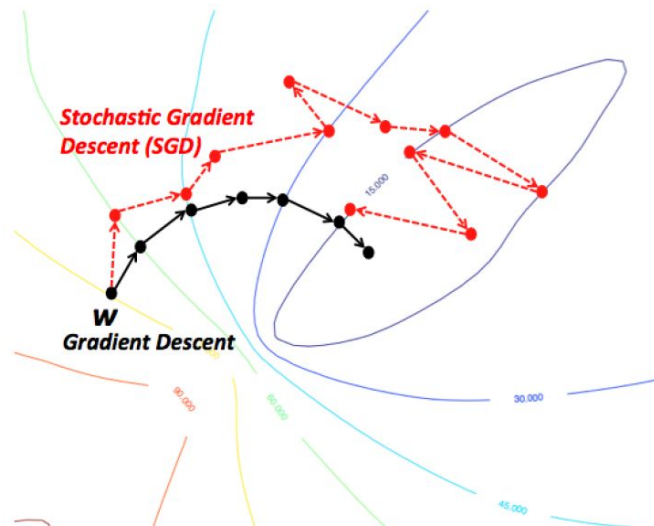


Stanford CS229 Lecture Notes by Andrew Ng



Stochastic Gradient Descent

- Calculate the gradient using one random data sample at a time
 - Takes many iterations to go over entire data set
 - Each time the full data set is covered is an "epoch"
- Makes for as noisier descent, which can be useful at times
 - See saddle point




http://www.bogotobogo.com/python/scikit-learn/scikit-learn_batch-gradient-descent-versus-stochastic-gradient-descent.php

Stochastic Gradient Descent (SGD)

Loop {
 for $i = 1$ *to* m {
 $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$
 }
}

Size of dataset




<https://upload.wikimedia.org/wikipedia/commons/f/f3/Stogra.png>

Mini-Batch SGD

Loop {
 for $i = 1$ *to* m {
 $\theta_j := \theta_j - \alpha \frac{1}{|b_i|} \sum_{x \in b_i} \frac{\partial}{\partial \theta_j} J(\theta_j, x)$
 }
}

Number of mini-batches

i th mini-batch



<http://cs231n.github.io/neural-networks-3/>



Momentum

- Helps SGD by reducing oscillation and focusing on the relevant direction to move
- Gives SGD a 'short term memory' by introducing a velocity term.
 - adds a fraction γ of the past time step's update vector to the current update vector



(a) SGD without momentum



(b) SGD with momentum

Update Rule:

$$v_t := \gamma v_{t-1} + \nabla_{\theta} J(\theta)$$

$$\theta := \theta - v_t$$

γ - the momentum term, usually set to 0.9

Adagrad

- Adapts the learning rate to each parameter
 - *Large updates* for *infrequent* features
 - *Small updates* for *frequent* features
- Relies on accumulating information throughout training
 - Which will eventually bring the learning to effectively zero :(

Update Rule:

$$\theta_i := \theta_i - \frac{\alpha}{\sqrt{G_{ii} + \epsilon}} \cdot \nabla_{\theta} J(\theta_i)$$

G_{ii} - the sum of the squares of the gradients w.r.t θ_i at current step.

Adam - Adaptive Moment Estimation

- Comparative to a combination of Momentum and Adagrad
- Exponentially decaying gradient sum for update - akin to momentum
- Gradient squared sum exponentially decays - solves adagrad problem.
- Gaining in popularity - generalizes well

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

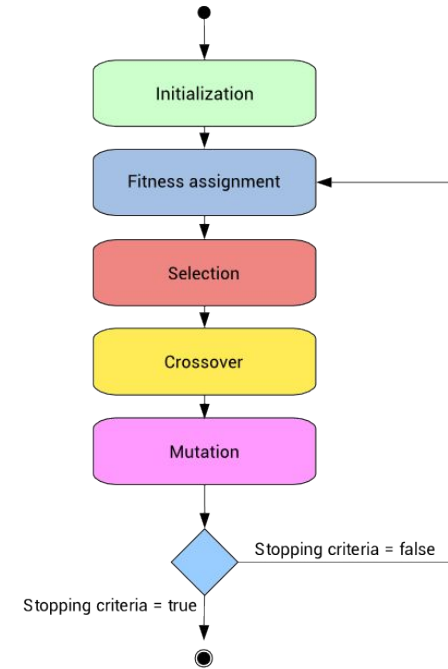
$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

Update Rule:

$$\theta_i := \theta_i - \frac{\alpha}{\sqrt{\hat{v}_t + \epsilon}} \hat{m}_t$$

Anything else besides Gradient Based Learning?

- Yes, but GD is fairly easy and effective
- Another cool option is using evolutionary models - essentially a guided random search
- Need three things for evolution
 - Variation
 - Heritability
 - Selection



https://www.neuraldesigner.com/blog/genetic_algorithms_for_feature_selection



Download the Jupyter Notebook:

<https://goo.gl/WZDwgC>


Gradient Descent for Linear Models

For the BUMIC gradient based learning workshop - fall 2018

Random Model

Let's generate some random data and make a random prediction for the best fit line. As you probably expecte

```
In [1]: 1 import matplotlib.pyplot as plt
        2 import numpy as np
        3
        4 # Generate some random linear data
        5 np.random.seed(42)
        6 X = np.arange(0, 10, 0.2) + np.random.normal(size=50)
        7 Y = np.arange(0, 10, 0.2) + np.random.normal(size=50)
        8
        9 # Hypothesize a line - just a random guess right now
       10 h_slope, h_intercept = np.random.rand(), np.random.rand()
       11
       12 # Create a list of values in the hypothesized line
       13 abline_values = [h_slope * x + h_intercept for x in X]
       14
       15 # View out hypothesis
       16 plt.scatter(X, Y)
       17 plt.plot(X, abline_values, 'r')
       18 plt.show()
```



References & Further Reading

1. Ruder, Sebastian. "[An overview of gradient descent optimization algorithms.](#)" arXiv:1609.04747 (2016).
2. Sun, Xu, et al. "[meProp: Sparsified Back Propagation for Accelerated Deep Learning with Reduced Overfitting.](#)" arXiv:1706.06197 (2017).
3. Kingma, Diederik, and Jimmy Ba. "[Adam: A method for stochastic optimization.](#)" arXiv:1412.6980 (2014).
4. Zeiler, Matthew D. "[ADADELTA: an adaptive learning rate method.](#)" arXiv:1212.5701 (2012).
5. Du, Simon S., et al. "[Gradient Descent Can Take Exponential Time to Escape Saddle Points.](#)" arXiv:1705.10412 (2017).
6. Dean, Jeffrey, et al. "[Large scale distributed deep networks.](#)" Advances in neural information processing systems. 2012.
7. Bottou, Léon. "[Curiously fast convergence of some stochastic gradient descent algorithms.](#)" Proceedings of the symposium on learning and data science, Paris. 2009.



Upcoming Events

- Introduction to Neural Networks
- First Hack Night on Sunday, September 30th
- First Paper Discussion within the next two weeks



Non-convex Optimization Issue

